

Machine Learning System Design:

➔ To prioritize your time spent

Eg Building a Spam Classifier

1) Find relevant features:

- a. Most frequently occurring n words in training set eg [10000-50000] words.

What options can we pursue to get the most efficient use out of our time?

- Collect More Data
 - o Eg Honeypot email accounts to get more spam mail examples.
- Develop sophisticated features based on email routing information (from email header)
- Develop Sophisticated features from message body
 - o 'discount' and 'discounts' considered the same word?
 - o 'Dealer' and 'deal'?
- Punctuation?

Workflow Approach:

- ➔ Start with a simple algorithm that is easily implemented, a Quick and Dirty Solution.
- ➔ Implement and test on CV data
- ➔ Plot a learning curve to evaluate if more data/more features, etc will help.
- ➔ Error Analysis

Error Analysis:

Manually examine the examples (In CV set) that the algorithm makes errors on.

- See if there are systematic trends that can be found in what kind of examples that it makes errors on
 - o Using such analysis, we can design new features based on these errors.

To test how effective these new features are, we conduct numerical evaluation:

Eg:

Should discount/discounts/discounted/discounting be treated the same?

- ➔ Use Stemming software / fuzzy string matching to categorise
 - o Note that such featurizing may also adversely impact the algorithm in other ways like University/Universe => clustered together even though they mean completely different ways.

Then evaluate the altered model against the initial model

w/o stemming: 5% error

with stemming: 3% error

Hence Yes it improves!

Distinguishing uppercase vs lowercase?

Without: 5%

With: 10%

No it does not.

- ⇒ Hence by seeing what works and what doesn't with a simple implementation we can decide what features are most important, and circumvent the problem of premature optimization.

Error Metrics:


Often times, the proportion of datasets may not be balanced
ie. there may be skewed classes

Eg: Logistic Regression Model on Cancer Diagnoses

Perhaps there is only 1% of actual positive test results in the dataset.

Precision/Recall

$y = 1$ in presence of rare class that we want to detect



Precision:

Percentage of true positives out of all predicted positives.

ie Specificity

Recall:

Percentage of True Positives out of actual number of positives

ie Sensitivity

We can adjust the sensitivity and recall by controlling the threshold value.

Eg Logistic Regression:

Threshold = 0.7

Predict 1 if $h_{\theta}(x) \geq 0.7$

Predict 0 if $h_{\theta}(x) < 0.7$

ie we only predict positive if we are confident more than the threshold about the prediction.

This causes the precision to increase, at a cost to the value of the recall.

Threshold = 0.3

Predict 1 if $h_{\theta}(x) \geq 0.3$

Predict 0 if $h_{\theta}(x) < 0.3$

This causes the recall to increase, at a cost to the value of the precision.

ROC curve,

The Precision and Recall can be plot against each other in a ROC curve.

F1 Score:

F1 Score is a single metric that helps measure the total efficacy of a model.

The formula:

$$F_{score} = 2 \frac{PR}{P + R}$$

Hence if $P = 1$ and $R = 1$, F score = 1

If $P = 1$, $R = 0$, F score = 0, same vice versa.

On Data:

Designing a high accuracy learning system.

Conditions for more Data as the solution:

Assuming Features $x \in \mathbb{R}^{n+1}$ has sufficient information to predict y accurately.

A useful heuristic test: Given input x , can a human expert confidently predict y ?

Use a learning algorithm with many parameters

Eg Logistic Regression/Linear Regression with many features / Neural Network with many hidden units => Low bias algorithm + Low variance

$J_{train}(\theta)$ is small

+ Use a large training set = Low chance of overfitting

$\Rightarrow J_{test}(\theta) \approx J_{train}(\theta)$

\Rightarrow Hence resulting in a effective model.